

Secure Scuttlebutt: An Identity-Centric Protocol for Subjective and Decentralized Applications

Dominic Tarr

ssb:@EMovhfFk4NihAKnRNhrf
RaqIhBv1Wj8pTxJNgvCCY=.ed25519

Aljoscha Meyer

TU Berlin, Germany
aljoscha.t.meyer@campus.tu-berlin.de

Erick Lavoie

McGill University, Montreal, Canada
erick.lavoie@mail.mcgill.ca

Christian Tschudin

University of Basel, Switzerland
christian.tschudin@unibas.ch

ABSTRACT

Secure Scuttlebutt (SSB) is a novel peer-to-peer event-sharing protocol and architecture for social apps. In this paper we describe SSB's features, its operations as well as the rationale behind the design. We also provide a comparison with Named Data Networking (NDN), an existing information-centric networking architecture, to motivate a larger exploration of the design space for information-centric networking primitives by formulating an identity-centric approach. We finally discuss SSB's limitations and evolution opportunities.

CCS CONCEPTS

• **Networks** → *Network architectures*; • **Software and its engineering** → *Publish-subscribe / event-based architectures*; • **Computer systems organization** → *Peer-to-peer architectures*; • **Information systems** → *Linked lists*.

KEYWORDS

Secure Scuttlebutt, Information-Centric Networking, Push vs. Pull

ACM Reference Format:

Dominic Tarr, Erick Lavoie, Aljoscha Meyer, and Christian Tschudin. 2019. Secure Scuttlebutt: An Identity-Centric Protocol for Subjective and Decentralized Applications. In *6th ACM Conference on Information-Centric Networking (ICN '19)*, September 24–26, 2019, Macao, China. ACM, New York, NY, USA, 11 pages. <https://doi.org/10.1145/3357150.3357396>

1 INTRODUCTION

A simple conceptual architecture for community applications consists of a global data pool to which every person can contribute and where every person can tap into the shared data – data sharing being the purpose of such applications. This model still is valid if one adds access control to the picture, either tied to the data (encryption giving access to content only to entitled holders of the decryption keys) or encrypting data in transit (login and TLS). Facebook and other centrally organized social app service providers fit well under

this global data pool model but have been strongly criticized for abusing their central provisioning position. The “decentralized web movement” [18] is the most visible technical response to this critique, pointing out implementation alternatives.

One of these alternatives is a project called Secure Scuttlebutt (SSB) that started in 2014. After several iterations of protocol design and implementation, SSB has become a stable service for over 10,000 users offering them rich media community applications with strong cryptographic protection (end-to-end encryption and metadata privacy) and running in pure peer-to-peer mode.

Selective Complete Log Replication

SSB relies on the core insight that each participant is only interested in a *subset* of the global data pool, thus it is feasible to locally store all the data a participant is interested in. To partition the data pool, all data is associated with the *identity* that produced it. Participants select their slice of the data pool by specifying the set of identities whose data they care about. This creates a “social graph” along whose edges data flows (Figure 1). Even as the overall system scales, the amount of data any single peer is interested in and thus needs to handle stays roughly the same.

Each participant can publish data to their single-writer, append-only log. This choice of data structure allows efficient replication and verification of the integrity of received data. Replicating these larger slices of the data pool comes with an unusual set of tradeoffs, discussed throughout the paper. As it turns out, replicated logs form a solid foundation for implementing many classes of applications.




Figure 1: SSB's “Internet of Identities” – Users A, B and C replicate logs (a, b, \dots) based on whom they follow: C does not follow A, hence has no log a . A and B follow each other such that when A follows C, A will get C's log c via B: new content is pushed directly if possible and through intermediary friends if necessary.

Subjective Reader

Because replication in SSB is selective and driven by a peer's social graph, different end devices will have access to different sets of log replicas, leading to different views of the world, which we call a



For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to the web page <http://creativecommons.org/licenses/by-sa/4.0/>

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

ICN '19, September 24–26, 2019, Macao, China

© 2019 Copyright held by the owner/author(s).

ACM ISBN 978-1-4503-6970-1/19/09.

<https://doi.org/10.1145/3357150.3357396>

“subjective reader” approach. SSB considers this a desirable property: each peer is free to consider data sources of its own choosing instead of having to feed from a centrally provisioned or otherwise converged view. While it is possible to implement consensus protocols over SSB, or to designate central data aggregators from which many peers consume the consolidated outputs, the SSB network itself deliberately doesn’t offer consensus services nor central content (directories etc). In Section 3.3 we will show the implications of this technological choice on the structure of distributed applications that can only read from and write to local logs.

Novelty

Putting complete replication of individual append-only logs at the core of SSB’s protocol avoids several hard problems in distributed systems. **First**, it is a radically decentralized approach requiring only minimal specification-level coordination among the participants but no run-time checks or configuration management. **Second**, although append-only data structures are well known for their benefits and are at the core of crypto currencies’ consensus finding, SSB uses logs without any consensus properties. The issue is deliberately sidestepped, but all necessary building blocks are provided to higher layers. **Third**, crafting a cryptographic ID system and maintaining a social graph that informs routing creates a very narrow filter: it implements a receiver-driven approach where data only flows where it is needed and provides flexibility in the actual data dissemination strategy. **Fourth**, it makes every peer a publisher by design. This property goes beyond the decentralized approaches like DAT [3] or IPFS [6] which assume that there exist replication servers but keep the separation between a data transport network and a server layer. **Last but not least**, log replication leads to a distributed system with inherent high resilience as any communicating element carries a persistent copy of the data. In traditional distributed systems, coordinating the data persistence as a basis for resilience is often an add-on task, or requires at least a special recovery service.

Comparing SSB to NDN

Despite SSB data replication being currently implemented as an application protocol (layer 7 in the OSI stack [8]), we believe that its underlying principles are worth studying from a network API perspective (layer 3 in the OSI stack) to highlight regions of the design space that could be further investigated. We sketch here some aspects in which SSB differs from *Named Data Networking* (NDN) [5], a popular proposal for Information-Centric Networking (ICN). A longer exposition of SSB’s underlying communication model is available in a separate publication [42].

In SSB, the delivery of information is organized around *named data streams* for signed events. The basic unit of addressing is a full log that might still produce new event messages in the future. The SSB streams guarantee *reliable causal ordering* [11] and *authenticity*.

Delivery of streams follows a *push model*: once a receiver has expressed interest in a stream, new items are transferred automatically without being requested individually. Flow-control (back-pressure) in the current overlay implementation is done implicitly by the TCP connections used to deliver data among peers.

Streams are tied to a single *identity* with a corresponding *public key*: only this identity can produce new items in the stream with the correct signature. Because the signature ensures the integrity of the items, they can be served from anywhere and by anybody. Moreover, users independently create the key that serves as an identity. Streams are therefore *self-certifying* [24] and their integrity does not rely on an external trust anchor nor a central naming authority or other name coordination and allocation mechanisms.

In contrast, NDN embodies quite different design choices. First, the basic elements of networking are *single pieces of data*, identified by hierarchical names similar to paths in a filesystem. Second, data is accessed via a *pull* model: a consumer issues an *interest* in a name, and the network delivers the corresponding data. Third, many NDN schemes rely on a hierarchical trust model to issue certificates that can in turn be used to sign individual pieces of data.¹

It is not clear that either design can subsume the other: one can implement SSB over NDN or the opposite but each option comes with significant runtime costs. Still, there could be an intermediate territory where a future synthesis of the two approaches may emerge. We therefore analyze the trade-offs that appear in various ways of layering NDN and SSB to enrich the discussion around future developments in ICN.

Structure of this paper

We start out by giving an overview of the SSB protocol in Section 2. Next we describe common patterns of how applications can be built on SSB (Section 3). We show how SSB relates to other networking protocols, first through a detailed comparison with NDN (Section 4) and then in the broader context of related work (Section 5). We conclude the paper with an outlook on some of the “work in progress” (Section 6) and an evaluation of the problems (Section 7) and benefits (Section 8) of the SSB approach.

2 SSB ARCHITECTURE AND PROTOCOL

In SSB, each user is identified by an ed25519 [7] keypair. Since anybody can generate a random keypair with very low probability of multiple peers generating the same keypair, no central authority is necessary for introducing users to the system. Conceptually, SSB is a network of identities that connect to each other (physical topology) and share mutual or unilateral interest in the other peer’s data (social graph), as shown in Figure 1. A node running the SSB protocol is called a *relay*. The identity that holds the private key of a log is called its *author*.

The *single-writer append-only logs* of SSB consist of entries (called *messages*) that include a *backlink* in the form of a cryptographic hash of the previous message (or a special indicator for the first message of a log). The most distinguishing feature of this linked list, when compared to a regular blockchain, is that each SSB user maintains their own log and cryptographically signs all their (and only their) messages. Messages whose backlink points to a message in a different log (i.e. by a different author) are considered invalid and will be rejected by SSB relays.

These constraints still allow creation of arbitrary trees rather than logs. To enforce log structure, each SSB relay checks that every

¹See [44] for a proposal that leverages a web of trust model in a decentralized chat application built on top of NDN.

message has exactly zero or one incoming backlinks. If it has more than one, the log is considered *forked*. All messages from the point of the fork onwards are ignored, the log cannot be appended to anymore.

Concretely, each message, which may not exceed 4 KB, contains the following pieces of data [28]:

- the *backlink* to the previous message, or a null value
- the public key of the message’s *author*
- the *sequence number* of the message, which must be one more than the sequence number of the previous message, or exactly one if it is the first message of the log
- a claimed *timestamp* of when the message was created
- a *hash* indicator that specifies the concrete hash function that was used to compute the backlink
- the *content* of the message
- the author’s *signature* over all the previous data

In the current version of SSB, *content* is a JSON object that must contain a *type* key that serves as a hint for how the content should be interpreted. SSB enforces that the content is valid JSON by rejecting any malformed message. Encrypted content is represented as a base64-encoded string, together with a tag that signifies which encryption algorithm was used.

SSB defines a format for encoding specifically the public keys of identities and the hashes of messages and blobs (see below) as strings. This allows applications to scan the content of messages from other authors for such references, e.g. in order to create database indices (see Sect. 3.3).

Replication

The principal function of SSB relays is to connect to other relays and exchange log *updates*. To do so, relays maintain a point-to-point encrypted [4] overlay network over which they run a gossip protocol. When two relays start gossiping, they exchange the current sequence numbers of all logs they are interested in. If a relay receives a lower sequence number for a log than it sent, it transmits the messages that the other relay is lacking. If at a later point a new message of the log becomes available to a relay, it is automatically *pushed* to the connected relays. As an optimization, this *eager* gossip is only performed over the edges of a spanning tree, which itself is maintained via the plumbtree [22] protocol. In classic peer-to-peer fashion, clients (leaf nodes) are no different than relays except that they usually include some graphical user interface and perform application logic.

In addition to this primary replication mechanism, SSB provides two other ways of exchanging information. *Blobs* are content-addressed pieces of free-form data, typically images or other documents larger than the 4KB limit, that are referenced from messages but are not part of any log. They are not widely disseminated automatically, but rather fetched on demand via a simple request-flooding protocol. *Out-of-order messages* are a similar mechanism to address and fetch messages on demand via their hashes.

SSB Relays as an Application Platform

Beyond replicating logs and checking the validity of update messages, an SSB relay offers an API to its peers. Peers can host arbitrary programs that issue remote procedure calls (RPCs) to the relay. The

exposed functionality includes appending to a log (if you know its private key), reading from logs, requesting which logs a relay should replicate, and fetching blobs and out-of-order messages.

The reference implementation of the SSB relay [40], written in JavaScript, also includes a mechanism for loading *plugins* into the relay to extend its functionality. There are a few default plugins: conceptually these can be thought of as client programs that are always running. Of particular importance are those that guide the replication process. The *friends* plugin [41] scans the relay’s log for specific messages that indicate which other identities the author *follows*. The plugin then instructs the relay to fetch and replicate these logs. These other logs might of course also contain some of these messages. The friends plugin transitively replicates these friends-of-a-friend logs as well, up to a configurable maximum distance in the friends graph.

Beyond “befriending” other authors through follow messages (i.e. messages of type follow), an SSB user can control the shape of their social graph via special block messages which limit the transitive log replication. Both the follow as well as block messages are overheard by relays, through scanning all received log updates, and inform them about where updates should be delivered (or not). Decisions about whom to replicate can be –and in the current system is– guided by the content of the very data that is replicated. By storing the relevant information inside the author’s log (as opposed to a local or central database), other peers can use this information to guide their decisions.

The overlay network also makes use of logs to store configuration information, in this case the SSB_ID-to-IP_address mapping: operators publish the static IP addresses of highly-available relays (called *pubs*) to their log. When an SSB relay needs to connect to the overlay, the responsible plugin can scan any locally available log replica for this information.

SSB’s Layered Architecture

It is worth noticing that SSB spans three independent layers of protocols. The most fundamental protocol is the message format: all peers need to agree on what constitutes identities, valid messages, and how to compute hashes to address messages and blobs. This is the “thin waist” of SSB (see figure 2).

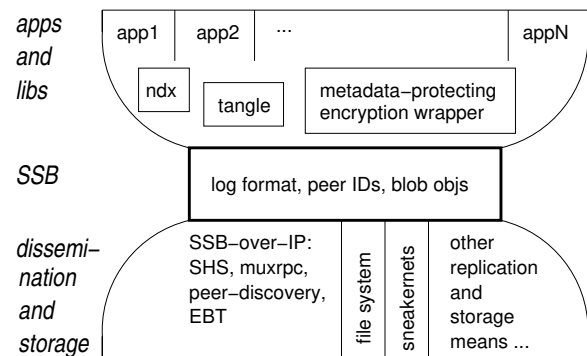


Figure 2: Secure Scuttlebutt’s protocol stack.

Next (below) is the specific mechanism by which relays exchange data. The default RPC mechanism is one option, but alternative

mechanisms such as distribution via a sneakernet could also be used. Different peers that do not share a common replication mechanism could still interact indirectly, as long as there are some relays that understand multiple replication protocols.

In other words, the core logical replication protocol by which a relay serves its clients is fully independent from the actual dissemination protocols. And finally, the publishing and interpretation of application data in such messages is again a separate affair that is layered on top of the thin waist.

3 DISTRIBUTED APPS AND DATA STRUCTURES OVER SSB

The replication model of SSB enables many collaborative applications to be implemented easily by abstracting much of the complexity in distributing the updates. However, implementing such applications still comes with challenges. To introduce them, we first discuss in detail the implementation of a user directory to introduce the implementation approach, then briefly cover other applications currently in use, and finally discuss some core issues that are shared between all applications.

3.1 Example: SSB’s user directory

‘about’ is SSB’s user database i.e., an application that associates cryptographic IDs with (typically) human-readable attributes. A single message format has been defined to this end:

```
'content': {
  'type' : 'about',
  'about' : about_id,
  attr_name : attr_value // multiple times
}
```

The about app scans all logs for messages of type ‘about’ and constructs a database as shown in Figure 3, retaining the most recent attribute assignment found. In this database, an about_id is associated with a list of key/value pairs which are prepended by the publishing author’s ID. It is left to the end user’s ‘about’ application to subjectively select which of these bindings should be displayed. Currently the name, textual description and image attributes are understood by most SSB user interfaces and are used to substitute or decorate the cryptographic ID. If about_id and author_id are identical, this means that an attribute was self-chosen and then is usually rendered with a higher preference over key/value pairs assigned by others.

about_id		key	val
about_id_1	author_id_A		...
	author_id_B		...
	...		
about_id_2	...		
...			

Figure 3: SSB’s user directory data structure (after extraction from the logs).

In terms of CRUD² actions, creation happens once a new SSB identity adds its own about message to its log; reading the user database

²Create, Read, Update, Delete

is performed on the above data structure; updates are expressed by adding an about message –regardless whether it relates to the author itself or to another identity– to one’s own (!) log and all peers updating their extracted database; deleting a user entry is not possible, at least not directly (one would have to block that user ID as well as all IDs which wrote an update for that user).

There can never be confusion about the sequence or scope of attribute assignments because they are ordered by the log (and thus in time) and kept separate, per author ID. Note also the presence of the “subjective reader” property: the content of a peer’s user database is dependent on their position in the social graph. The “subjective” mindset is also visible by letting every user assign attributes to anybody, leaving it to the user interface (and human viewer) to select which of the self-chosen or given display names and images is most suitable for a given ID.

3.2 Profiles of other selected SSB apps

Multiple applications have been written by contributors and are used daily by the SSB community. The following selected examples represent alternatives to well-known services and they illustrate both opportunities and challenges of communication through replicated append-only logs.

Git-ssb [21] is an alternative to GitHub [2] that replicates git-based version-controlled code repositories through contributors logs. It provides an encoding of repositories in SSB logs, a bridge to interoperate with git repositories, and a web-based viewer to browse repositories. The object model of Git [14] has a similar structure to SSB’s logs. Other git operations, such as creating repository, are all SSB messages. Since any user can independently update the *same* repository, as defined by its creation message, consensus on the “official” master branch and its latest commit is enforced through social coordination. In case of concurrent updates to the same branch in the same repository [32], referencing both concurrent updates in a later merge commit in effect resolves the ambiguity.

Ssb-chess [23] is a correspondence chess application in which players can invite one another to play, alternatively share their next move until the game ends, and external observers can comment on the game. Because the rules of chess preclude concurrency, i.e. at any time there is always only one of the two participants that is permitted to modify the state of the chess board, a game can easily be represented as a linked list with nodes representing chess moves alternating between the two participants’ logs. Moreover only the participants, explicitly mentioned in the original invitation, are allowed to modify the state of the game. The implementation does not require concurrency management and is therefore conceptually straight-forward.

Gatherings [16] are alternatives to Meetup [1] that enables participants to signal their intention to attend or not attend to physical events. A gathering is defined by its creation message but otherwise has no fixed properties. Anyone that has a reference to the creation message may change its properties, such as location, start and end dates, description, and image, by publishing an update message. The value of those properties are the most recent set by anyone. Initially, recency was determined by the time of creation, as reported by the user’s client implementation (*self-stated creation time*). To be more robust to potential invalid timestamps however, some client

implementations have started using the time at which message updates are *received*, then disambiguate using the self-stated creation time.

3.3 Running Distributed Applications over Replicated Logs

“Infrastructure-less” distributed application as presented above become possible because central servers can be fully replaced by each peer working on its local set of replicated logs. In this subsection we discuss the particularity of this approach and its constraints.

A common pattern of SSB’s applications is that they heavily rely on local database support for organizing the data contained in the logs. Typically a map-reduce strategy is used where the map phase filters the logs and the reduce actions computes the latest application state.

In the user directory application (Sect. 3.1), the filtering is done by selecting only about messages for a specific target ID and the reduce action consists in accumulating the latest key-value pairs such that a more recent key-value pair replaces an older one if it was signed by the same `author_id`. The size of the replicated logs, although locally stored, would lead to very long response times if the map-reduce would be executed at render-time. Instead, almost each application will build indexes and aggressively cache state that was already aggregated. Should the indexes ever become corrupted (e.g. because the user interface app crashed in the middle of a complex indexing step), they can be fully regenerated from scratch.

An important aspect is whether the reduction step can be done in an incremental fashion by reusing previously computed application state. For example, counting the “likes” that a post receives works fine: incoming log extensions are indexed and if they are of the like type, the counter corresponding to the referenced message is incremented.

Other applications, however, may need a *full* re-evaluation of the reduce function each time the underlying index changes. An example for this case is a chat in form of sequence of post messages: if some identity is added to the set of followed identities, its log is incrementally replicated, and so are the posts of this identity. For each incoming new post, which may have been written very long ago, one has to insert it at the right place. This problem is shown in Figure 4 where a message has not yet been replicated to a client and, once it arrives, has to be properly inserted into the application-level data structure.

A simple solution (adopted in some SSB client software) is to use the timestamp claimed by the author of the post, and in this case one can reuse the existing time-sorted list and insert the new post. However, because an author could lie about the timestamp, the reduce function should do a topological sort based on the causality relationship with other posts and their replies, which form a directed acyclic graph.³ Insertion into the dependency graph may or may not lead to having to rerun the sort on the whole graph of postings. Clearly, the lack of a central server hosting the reference list of posts and being able to record a post’s submission time, leads

³Writers can facilitate the causal ordering by explicitly referring to the most recent messages in other logs they were aware of at the time of writing, regardless of whether the content was directly relevant to their own message. We call the graphs resulting from this pattern “tangles”.




Figure 4: Example of extracting application data spread over multiple collaborating logs and dealing with not-yet-delivered data.

to more complex client software that must prepare for and defend against a broad range of adversarial data found in the logs.

3.4 Synchronization and Eventual Consistency

SSB’s basic log replication service synchronizes peers in a consistent way: due to the hash chaining, events (represented by messages in a log) will be delivered in the order they happened and replica content will be consistent. This does not instantly lead to consistent shared data structures, though, if the corresponding events are spread over multiple logs. Instead, the natural guarantee is that of *partial eventual consistency* where all peers will see the same reduced application state *if they share the same log set* after sufficient replication progress. Because of the eager, push-based forwarding of messages, consistency can be reached quickly, even if there is no end-to-end connectivity.

Eventual consistency is the hallmark of Conflict-free Replicated Data Types (CRDTs, see [38]) which are directly applicable to the SSB setting as they only assume a reliable and (sometimes) in-order delivery of update messages. Potentially, CRDTs permit to implement global data structures featuring eventual consistency without coordination effort (thus are fully scalable). The caveat here is that SSB peers do not necessarily see all involved logs due to their position in the social graph which controls replication wherefore consistency is always modulo that fact. For example, like counts will be eventually consistent with respect to the same set of followed authors but not globally, at least if they are directly counted. Other applications relying on reduction via set union may learn from state that stems from beyond the circle of followed authors.

More research is needed to understand the constraints brought by the combination of coordination-less interaction with partial log replication, but SSB’s rich set of applications used on a daily basis is an encouraging sign that eventual consistency in combination with subjective replication is a “good enough” basis for real decentralized services.

4 COMPARING SSB WITH NAMED DATA NETWORKING (NDN)

In this section, after a brief introduction to NDN we compare and relate SSB to NDN in three different ways: layering SSB on top of NDN, layering NDN on top of SSB, and a hybrid mode that combines

features of both. The purpose is to shed light on the sometimes implicit assumptions behind SSB and NDN and show a larger design space for future ICN developments.

4.1 Named Data Networking

NDN is an evolution of the Content-Centric Networking proposal that was publicly presented by Van Jacobson in 2006 [19]. Both aim to address shortcomings of using the Internet Protocol (IP) [13] for data dissemination from a single source to a number of users. In the IP, the routing layer only deals with the delivery of data packets from a source to a destination, regardless of their content: when multiple users request the same content from different machines, the routing layer therefore has to deal with redundant data transfers and is prone to content tampering by intermediate routing nodes. Some of the major goals of NDN are therefore to optimize data distribution for large content providers while guaranteeing the integrity of content. NDN currently achieves those aims by: (1) initiating data transfers *after* the interested users are known by the network (*pull-model*), (2) leveraging *existing certificate infrastructure* to authenticate the content, and (3) *naming* individual pieces of data, using a naming scheme that reflects the hierarchical organization of major content providers, such as universities, governments, and major media companies.

Technically, a receiver has to request content by name –in a so called Interest packet– and at most one matching content is returned in a corresponding Data packet. The Data packet includes the content’s name and is signed such that a forwarding node can verify the correctness of the name-to-content binding.

```
--> I('/ndn/some/item')
<-- D('/ndn/some/item', data, signature)
```

Checking the validity of a signature requires additional certification data which a forwarding node can fetch using the standard Interest/Data packet pattern. Validated data packets are typically cached such that subsequent requests for the same name can be served from in-network memory.

Routing rules are based on name prefixes, which aggregates all data items made available by a publisher. In a forwarder node, incoming Interest packets are matched against these prefixes on a longest-prefix matching basis, yielding the interface(s) to where an Interest has to be forwarded. Interests for the same name that arrive close in time are deduplicated using a PIT (pending interest table). On the return path, a data packet is copied to all interfaces from where a corresponding Interest came in, and the PIT entry is deleted.

4.2 Comparison with SSB

Similar to NDN, SSB organizes distribution around *content*, instead of the *machines* that are interacting. In contrast to NDN, the design aims at *individual users as publishers* instead of larger organizations with the following technical consequences. SSB addresses a *stream of data* tied to a particular identity instead of individual data items. SSB *eagerly broadcasts* content as soon as possible (*push-model*), leveraging the abundant storage available in peer’s devices for replication. The replicas are determined based on social interests between *peers*, achieving a similar aim as the Interest packet but once for an entire *stream* of data and with persistence by default,

rather than for each individual items with temporary caching. In SSB, there is no distinction between consumer and producer roles: every client must be able to produce (signed) log entries. SSB avoids the use of certificate authorities to secure the respective signing keys, instead relying on self-generated cryptographic key pairs and *trust* between peers established over repeated social interactions to establish credibility in a particular identity.

The different application context of NDN and SSB has an impact on ID management. In NDN, users (content consumers) are anonymous and their interest in the same content can be aggregated if it comes from shared routes. In SSB however, recipients also have an ID with an associated log in which they declare their interest in another ID. Said differently, NDN works with repo IDs (prefixes) on top of which we have IDs for content (= content names extending a repo ID). In NDN, IDs have no role for the receiver or in the replication process except that forwarding validates the origin of data items. On the other hand, these repo IDs must be globally routable through some unspecified routing protocol outside the NDN specs. SSB also has producer-side IDs, but it is mandatory that clients also have an ID because otherwise they could not publish their replication needs (towards SSB’s routing logic).

The differences in design decisions make the combination of SSB and NDN hard to efficiently layer one way or the other, as shown in the next sections and illustrated in Figure 5.

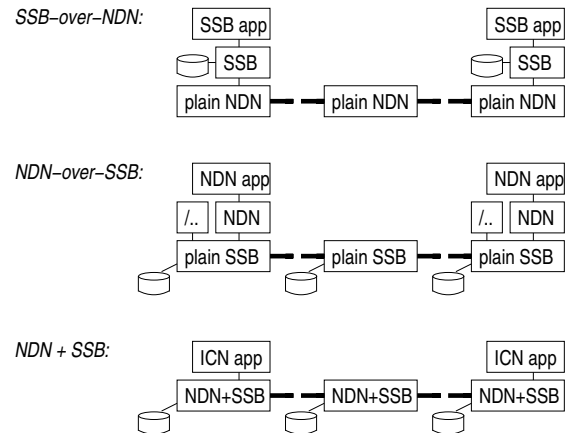


Figure 5: Three different layerings of SSB and NDN.

4.3 SSB over NDN

Emulating SSB over NDN means emulating a push-based system over a pull-based one, and an identity-centric system over a name-based one. Both turn out to be problematic.

Implementing push with the pure request-reply model of NDN comes down to two basic options [12]: the producer could send an Interest to the consumer, to signal that the consumer should itself issue an interest in the newly available data. This approach incurs a high latency penalty and leaves a trail of in-network state.

The other approach is regular (per-item) polling: the consumer periodically signals interest for some data the producer may or may not have created yet. In its simplest form, this can be done by publishing data under a name that ends with a sequence number

that is incremented with each produced piece of data. Under this model, the consumer can decide how many items into the “future” to poll for simultaneously. This whole process can be abstracted over with a consumer-side library [29, 35].

With the current design of NDN, polling is resource intensive. A natural extension of per-item polling is the inclusion of long-lived or persistent Interests [30]. But even then, a pull implementation would be inefficient: polling ahead for multiple items effectively amounts to controlling back-pressure through a sliding window, comparable to TCP. But unlike TCP, this window could only be manipulated by one item per (interest) packet. Introducing some form of sequence number arithmetic to increase efficiency would necessitate to drop the concept of purely opaque names.

Implementing the pull aspects of SSB over NDN would therefore cost either time (interests triggering interests), space (polling), or it would require significant changes to NDN (long-lived interests + non-opaque names) that would effectively bend NDN towards a “name-based SSB”. We will explore this option in more depth section 4.5.

4.4 NDN over SSB

While the current design of NDN is not well suited for implementing SSB, implementing NDN over SSB would also be inefficient. SSB would have to implement three distinct NDN features: the hierarchical name space, the pull-model and NDN’s trust system.

Again assuming rough equivalence of NDN data packets and SSB messages, i.e. a triple $\langle \text{name}, \text{content}, \text{signature} \rangle$, the pull-model part is easy to answer: either some item is already in one of the eagerly replicated local logs, or it is not available yet (because SSB is push-based).

The major problem is NDN’s hierarchical namespace which is a globally shared construct with the service level agreement that any (existing) item referenced through this tree can be fetched. Even if a delegation model is used, this global resource introduces a central authority, or at least a consensus algorithm, to allocate prefixes. This entity would have a “well-known” SSB id and a log from where the rest of the SSB world would inform itself about its decisions. Once these prefixes are handled (by a special SSB app for supporting NDN’s namespace, depicted as $\boxed{/. .}$ in the third subfigure of Fig. 5), repo IDs can be introduced such that an end device can address them and request content from them. However, in SSB’s worldview, a repo would have to follow all potential customers i.e., to learn about their IDs, otherwise these customers cannot express interest in some content (which could be delivered through log replication of transient SSB IDs, for example). When looking at NDN from the viewpoint of SSB, we realize that NDN has a social contract along the interest path: an NDN forwarder accepts any downstream node as a friend, accepts its interest packets (= pushed replica of the requests), and then relies on a similar contract with its upstream node. Following this insight, our NDN emulation would have to introduce “NDN forwarding providers” at SSB level. Once these “NFPs” are in place, we would also let them implement NDN’s trust model by validating content through NDN’s certificate authorities.

While it doesn’t seem strictly impossible to continue that emulation argument, it is already obvious from the above discussion that one would not benefit from SSB’s social graph replication mindset.

4.5 Combining NDN and SSB

While in the previous sections we tried to layer the two ICN network architectures, NDN and SSB, on top of each other to only unsatisfactory results, we explore the potential of bringing only a subset of SSB’s functionality into NDN in this section. The focus is on SSB’s append-only log which enables a “controlled push” service, answering an often-heard request towards NDN from distributed systems developers. Other elements of SSB like the symmetry between producers and consumers, or the use of the social graph to inform the routing, are left out.

A core critique against network-level push is that a producer can flood all of its consumers: letting consumers ask for one individual data item with an interest packet effectively blocks any flooding attack because the PIT state is consumed by the first data packet that satisfies the request. We sketch - and suggest to explore - a system where consumers subscribe to some data stream *and* keep a way to apply backpressure. To this end, we map SSB’s log replication semantics to an NDN-style, packet-level protocol design.

In this protocol design, data is organized in *feeds* (comparable to SSB logs) that consist of individual data items (corresponding to SSB messages) that can be requested individually via traditional NDN-pull using the feed’s name plus the item’s sequence number. In addition, there is a special *long-lived interest* (LLI) mechanism that allows a consumer to subscribe to a feed for a limited number of packets. An LLI is a tuple of a feed name, a sequence number, and a *send credit*. The send credit tells the upstream node how many items following the sequence number are needed to consume the PIT entry, effectively defining a (finite) subrange of a feed.

As soon as new items are available at an upstream node, the data is pushed towards all subscribing downstream nodes, up to their expressed send credit. Forwarding nodes store (cache) the items of a feed and check that new events correctly extend the local copy of the feed.⁴

Lost packets are detected when a received sequence number does not match the expected one. Different retransmission strategies could be employed, for example selective retransmission via a classic NDN Interest, or go-back-N with an LLI packet for a range that starts at the first missed packet. Compared to an emulated PUSH by polling via multiple Interests in parallel, less PIT state is needed for such a go-back-N strategy: the feed concept allows ranges to compactly encode multiple names. This is an example of an optimization that becomes possible because the network knows about the structure of feeds. Investigating other efficient retransmission mechanisms deserves further research.

We call this approach a *controlled push* for two reasons. First, consumers guard against buffer overrun by expressing a send credit (which forwarding nodes can aggregate and adjust according to the available cache memory): in case of setting all send credit values to one we get NDN’s classic pull. Second, the cryptographically secured append-only logs impose a single feed source: the network benefits from this information because only a limited number of the most recent items need to be cached to serve most consumers.

⁴ Asserting feed integrity involves more than just checking an item’s signature: the new item’s sequence number must be in order, the backlink must correctly reference the previous item, and the feed should not be forked. A nontrivial problem is how nodes should react to forked feeds: ideally, the information that a fork occurred should be propagated. We leave this open for future work.

The resulting “flow-controlled reliable stream” network-level service is attractive for ICN application writers in several ways: once requested, asynchronous data is pushed with zero delay, all events are ordered and are checked for their integrity (single-source feed), and consumers stay anonymous and don’t need their own routable prefix. Additional benefits exist due to the network being aware of the feed semantics, for example being able to request in-network retransmission when observing missing feed entries, instead of having to wait for the PIT entry to time-out.

While unfortunately some properties of SSB are left aside during such an import of push semantics into NDN, we see interesting research opportunities along the sketched path where a data-structure aware network substrate can provide advanced and classically dangerous services like push *because* of the data structure’s property. To us, the price to pay in terms of PIT state seems minimal, and justified for gaining native push semantics in NDN.

5 RELATED WORK

Some basic ideas behind SSB can be traced back to the nineties, like for example *secure logging* [36] and *secure relative time-stamping* [17]. The major innovation of SSB is to use these techniques for disseminating data through a gossip protocol in a network of untrusted peers, effectively implementing a push-based information-centric network. SSB’s messages, named by $\langle id : seqno \rangle$, are a special case of DONA’s naming schema [20], where sequence numbers can be regarded as totally ordered labels.

SSB’s push-based content dissemination approach is also underlying middle-ware systems like *Linda* [15]. Linda offers a global data pool abstraction where distributed processes can store and consume objects without locality references: the effects of a `wr()` operation are propagated automatically such that processes being blocked on a `rd()` could be resumed immediately.

In the area of delay-tolerant networking, systems like *HAGGLE* [37] and *SCAMPI* [34] also aim to leverage social dynamics between users. These systems correlate social proximity with physical network connectivity to enhance performance and availability of applications. SSB primarily uses social dynamics to determine *which* data should flow to *whom*, not *how* it should flow to *where*. In SSB, identities and their relations are first-class, whereas HAGGLE and SCAMPI rely on inferring them by extensively monitoring user activity.

The use of logs itself has a long tradition in distributed systems, especially in operating systems (*write-ahead logs* in journaling file systems) as well as distributed databases. More recently, in the cloud context, resilient event ordering protocols like *RAFT* [33] have been proposed that also rely on replicated logs. Although logs are used at various places of distributed systems, this data structure is typically not exposed to the communicating parties, while SSB exactly rests on letting apps interact directly with the secured single-author logs.

Selective Hearing [26] uses a gossip protocol to disseminate monotonically growing sets of updates to provide a runtime to the Lasp [25] programming language. The general architecture is similar to that of SSB, the most striking difference is that Lasp is by design restricted to CRDTs. SSB can be considered more low-level, developers are free to choose a strategy for dealing with concurrency and eventual consistency. Selective hearing was developed

in a more traditional research approach, so it glosses over some of the difficult problems encountered in the “real world” such as user onboarding and byzantine peers. Their “practical large scale evaluation” [27] consists of 1000 well-controlled nodes running a toy application in the cloud, whereas SSB with its roughly 10000 users is more battle-proven.

6 FUTURE WORK

So far we have mostly restricted our presentation to those features that are implemented today as part of SSB. In this section we will describe further extensions, namely a potential revision of SSB’s log format and the operational challenges for scaling SSB beyond its current user base.

6.1 Partial Replication

By using a linked list of messages as the underlying datastructure for a log, a message can only be verified to be a valid element of a specific log in time linear to its sequence number. Since all previous messages need to be available for verification, this also implies linear storage overhead. More sophisticated datastructures could reduce this to logarithmic overhead, both anti-monotone binary graphs [9] and threaded authentication trees [10] would be suitable and would only require a single additional hash per message.

An interesting problem in this context is how peers would indicate the subsets of a log they want to receive. Specifying individual sequence numbers works fine, but degrades to a pull-based system. Instead semantic criteria are needed, for example subscribing to only messages of certain types. Finding a general framework for specifying partial subscriptions based on semantic frameworks is an interesting task. Care must be taken that malicious peers cannot silently suppress data that matches a partial subscription, this could be done by adding additional sequence numbers for each criterium.

6.2 Cryptographic Agility

SSB relies on multiple cryptographic primitives (signatures and hashes for the log format, encryption for the replication protocol): best practice mandates that cryptographic agility is supported [31]. All hashes and signatures in the logs include an indicator of the cryptographic primitive that has been used. At least in theory this means that the SSB protocol can introduce the use of new primitives as old ones become broken.

In this context, an open problem is how old messages can be “saved” once their signature data or hash references become insecure. The naive approach of republishing old messages with a new key changes the hashes of all those messages, thus breaking inter-message references. A similar discussion (and proposed solution) for NDN can be found in [43].

6.3 Multi-Device Support

If two different devices used the same SSB identity to publish messages concurrently, this would result in two competing hash chains with the consequence that peer relays would stop propagating at least one, if not both log extensions. It is thus recommended to create a distinct keypair per device. But this leads to bad user experience, such as having to follow or block identities on all device.

This could be mitigated by developing schemes that allow sharing the same private key across multiple devices to allow read-access, while enforcing mutual exclusion on writes.

A different angle is to write applications in a way that anticipates that there might be a one-to-many mapping from users to SSB identities. Since the messages in a single log are totally ordered but messages across multiple logs might only be partially ordered, it is not sufficient to naively treat a set of logs as a compound log. Instead, applications need to be designed from the ground up to deal with partially ordered sets of messages.

6.4 Replication Improvements

The currently used gossip-based default replication protocol does not protect against malicious intent such as for example eclipse attacks [39]. But whereas it is difficult to defend against these attacks in general, SSB can make use of data such as the friend graph to protect against them. A *follow* message can be interpreted as an expression of trust. Keeping a certain number of trusted peers in the views of the peer sampling service could protect against eclipse attacks.

Another area where the replication protocol could be improved is by using private set intersection when determining the set of logs that both parties are interested in. That way, untrusted peers would not be able to learn about new ids purely from the replication layer. Combined with an access control mechanism that only forwards data to authorized identities, this would provide resilience against bots “spidering” the network.

7 SSB CHALLENGES

In this section we critically review limitations and challenges faced by identity-centric systems such as SSB. We omit those problems that apply to SSB in its current state but that would be solved by the extensions presented in the previous section.

7.1 Privacy

SSB is an inherently pseudonymous system: anonymity is fundamentally incompatible with identity-centric message propagation. Furthermore, the architecture discourages ephemeral pseudonyms, favoring the creation of a rather stable network of trust to guide replication. Since all messages are signed, they are not refutable. Finally, all messages are immutable.

The cocktail of pseudonymity, non-refutability and immutability can be a serious risk to users. Personal details could fuel harassment, political statements could justify persecution, all data could serve as the basis of (future) discrimination. The risks can be reduced by taking care that pseudonyms cannot be traced to physical identity, compartmentalizing pseudonyms, using encryption, and only giving the messages to trusted parties. Still, participation currently favors privileged users for whom privacy issues are not critical. Consequently, applications must clearly inform users about the peculiarities of the virtual space they participate in to ensure users don’t share information that might be detrimental to them later.

7.2 Onboarding

Data can only be propagated to relays that specifically ask for it. When a new identity joins the system, it can only participate

effectively once someone subscribes to its log. The SSB community approaches this “onboarding” with multiple techniques. Pubs, acting as quasi-permanent SSB relays, can issue *invite codes* out-of-band. When a new user sends such a code to the pub, together with its self-generated public key, the pub automatically follows the user, requesting their messages in the process. As an additional onboarding mechanism, the reference relay implementation uses LAN multicast to discover nearby peers. This allows local onboarding where an established user can follow a new user in the same LAN.

7.3 Coordination

The *type* field of SSB messages can be regarded as a global resource without any central coordination regarding its usage. In the worst case, this can lead to multiple applications using the same *type* but in incompatible ways. Namespacing and random types reduce but don’t eliminate this problem.

Non-interoperable *types* are a very tangible symptom of a broader theme, the *plurality* of interpretations of message contents. Taken to an extreme, “the” community of SSB users could splinter into a multitude of mutually non-understanding fractions that use different kinds of messages or interpretations thereof. Supporting divergence can also be considered a feature because it mirrors the informal evolution of human languages over time, a property that is often overlooked or actively shunned in more centralized designs.

8 BENEFITS

Here we summarize some desirable properties of SSB. These go beyond the obvious productivity gains for application developers who don’t need to implement encryption, authentication and synchronization, as well as the automatic replication of the content through SSB’s push approach.

8.1 Resilience

The design of SSB intentionally avoids “global singletons”, or centralization aspects requiring consensus. SSB can therefore be characterized as a “collection of decentralized systems” that overlap to varying degrees. It is consequently highly resilient to failures, whether due to attacks or errors in the code base or in operations. Also, users do not need to depend on any single, privileged central authority, including cloud-based service providers and instead participate in deliberately isolated networks, for example limited to a family, or a specific local area network.

Since SSB applications only interact with the local replicas of logs, complete offline operation is automatically built in. Offline operation is simply a special case of a (temporary) network partition. Because this case occurs so often, the protocol is geared towards handling network partitions gracefully, further contributing to the resilience of SSB. In particular, all operations are delay tolerant.

8.2 Efficiency

Due to the “subjective reader” approach, all SSB relays and application programs can operate concurrently. There is no need for synchronization across relays, avoiding the overhead this might incur, and applications can fully embrace the properties of the

append-only logs: The monotonically growing logs are well-suited for implementing CRDTs and similar techniques.

By *leveraging* existing trust behind social interactions, instead of trying to eliminate it, as e.g. blockchain systems which establish consensus over trustless nodes do, proof-of-work and other computationally-intensive techniques can be replaced by *social signals* encoded as events in a log. The delay tolerance allows routing layers that can optimize for different tradeoffs, for example by minimizing the bandwidth required to disseminate updates rather than minimizing latency. Pushed further, the same approach could lead to infrastructure that is quite efficient in its usage of memory, bandwidth, and energy, making the overall required infrastructure sustainable with less resources than other approaches requiring always available, high-throughput, routing infrastructure. Leveraging existing social trust between participants can therefore provide clear technical benefits.

8.3 Plurality and Disintermediation

The freedom of applications to interpret data in whatever way they see fit increases the agency of users and application writers, to choose how to leverage the data they produce and for what purposes. SSB supports plurality also as a deliberate strategy to drive evolution. This also fosters the sharing of data between applications. For example, the ‘about’ information (Section 3.1) can be reused by all programs, freeing implementors from duplicating work and creating a coherent user experience across apps.

As another consequence of the subjective interpretation of data, there is no need for central coordination to introduce or evolve features: new uses can evolve based on the immediate needs of participants and then spread if the needs are more widely shared. Applications can simply start producing new kinds of messages, and interoperability works out with all users who share the same interpretation of those messages. This results in an organic evolution of features, without “the system” ever shutting down.

9 CONCLUSIONS

We presented Secure Scuttlebutt, a fully decentralized, peer-to-peer event-sharing protocol. The core novelty is that data replication occurs at the granularity of complete, self-certifying append-only logs of messages by a particular author. This approach leads to a simple, yet efficient information-centric service abstraction that lends itself well to a large class of applications. By embracing push-based eventual delivery and subjective interpretation of data, SSB gets to sidestep common sources of complexity. A community of multiple thousand users interacting through a variety of different applications confirms the viability of the approach.

The comparison with NDN shows that SSB’s paradigm of push-based, identity-centric data transfer comes with a different set of tradeoffs than NDN’s choice of pull-based, name-centric data transfer. Focussing on identities leads to challenges with respect to user privacy, but it also enables elegant, decentralized solutions to common problems with information-centric systems. Whether in the context of SSB or more generally, we believe that further study of identity-centric systems will lead to valuable insights and designs.

REFERENCES

- [1] 2002 - 2019. *meetup.com*. <https://www.meetup.com/>
- [2] 2008 - 2019. *GitHub*. <https://github.com>
- [3] 2014. *DAT Project*. <https://datproject.org/>
- [4] 2015. *Designing a Secret Handshake: Key Exchange as a Capability System*. <http://dominictarr.github.io/secret-handshake-paper/shs.pdf>
- [5] Bengt Ahlgren, Christian Dannewitz, Claudio Imbrenda, Dirk Kutscher, and Borje Ohlman. 2012. A Survey of Information-Centric Networking. *IEEE Communications Magazine* 50, 7 (2012), 26–36.
- [6] Juan Benet. 2014. Ipfs-content addressed, versioned, p2p file system. *arXiv preprint arXiv:1407.3561* (2014).
- [7] Daniel J Bernstein, Niels Duif, Tanja Lange, Peter Schwabe, and Bo-Yin Yang. 2012. High-speed high-security signatures. *Journal of Cryptographic Engineering* 2, 2 (2012), 77–89.
- [8] Neil Briscoe. 2000. Understanding the OSI 7-layer model. *PC Network Advisor* 120, 2 (2000).
- [9] Ahto Buldas and Peeter Laud. 1998. New linking schemes for digital time-stamping. In *ICISC*, Vol. 98. 3–14.
- [10] Ahto Buldas, Helger Lipmaa, and Berry Schoenmakers. 2000. Optimally efficient accountable time-stamping. In *International Workshop on Public Key Cryptography*. Springer, 293–305.
- [11] Christian Cachin, Rachid Guerraoui, and Luís Rodrigues. 2011. *Introduction to reliable and secure distributed programming*. Springer Science & Business Media.
- [12] Antonio Carzaniga, Michele Papalini, and Alexander L. Wolf. 2011. Content-based Publish/Subscribe Networking and Information-centric Networking. In *Proceedings of the ACM SIGCOMM Workshop on Information-centric Networking (ICN '11)*. ACM, New York, NY, USA, 56–61. <https://doi.org/10.1145/2018584.2018599>
- [13] Vinton Cerf and Robert Kahn. 1974. A Protocol for Packet Network Intercommunication. *IEEE Transactions on communications* 22, 5 (1974), 637–648.
- [14] Scott Chacon and Ben Straub. 2014. *Pro git (2nd Edition)*. Apress. <https://git-scm.com/book/en/v2>
- [15] David Gelernter. 1985. Generative Communication in Linda. *ACM Trans. Program. Lang. Syst.* 7, 1 (Jan. 1985), 80–112. <http://doi.acm.org/10.1145/2363.2433>
- [16] Piet Geursen. 2017. *patch-gatherings*. <https://github.com/pietgeursen/patch-gatherings>
- [17] Stuart Haber and W Scott Stornetta. 1990. How to time-stamp a digital document. In *Conference on the Theory and Application of Cryptography*. Springer, 437–455.
- [18] Internet Archive. 2018. *Decentralized Web Summit 2018, Jul 31 – Aug 2, San Francisco*. <https://decentralizedweb.net/>
- [19] Van Jacobson. 2006. *A New Way to look at Networking*. <https://www.youtube.com/watch?v=oCZMoY3q2uM>
- [20] Teemu Koponen, Mohit Chawla, Byung-Gon Chun, Andrey Ermolinskiy, Kye Hyun Kim, Scott Shenker, and Ion Stoica. 2007. A Data-oriented (and Beyond) Network Architecture. In *Proceedings of the 2007 Conference on Applications, Technologies, Architectures, and Protocols for Computer Communications (SIGCOMM '07)*. 181–192. <https://doi.org/10.1145/1282380.1282402>
- [21] Charles Lehner. 2018. *Git-SSB: Social Coding on Secure Scuttlebutt*. <https://git.scuttlebot.io/%25n92DiQh7ietE%2BR%2BX%2F1403LQoyf2DtR3WQfCkDKlheQU%3D.sha256>
- [22] Joao Leita, Jose Pereira, and Luis Rodrigues. 2007. Epidemic broadcast trees. In *2007 26th IEEE International Symposium on Reliable Distributed Systems (SRDS 2007)*. IEEE, 301–310.
- [23] Gordon Martin. 2017. *ssb-chess*. <https://github.com/Happy0/ssb-chess>
- [24] David Mazières and M. Frans Kaashoek. 1998. Escaping the Evils of Centralized Control with Self-certifying Pathnames. In *Proceedings of the 8th ACM SIGOPS European Workshop on Support for Composing Distributed Applications (EW 8)*. ACM, New York, NY, USA, 118–125. <https://doi.org/10.1145/319195.319213>
- [25] Christopher Meiklejohn and Peter Van Roy. 2015. Lasp: A language for distributed, coordination-free programming. In *Proceedings of the 17th International Symposium on Principles and Practice of Declarative Programming*. ACM, 184–195.
- [26] Christopher Meiklejohn and Peter Van Roy. 2015. Selective hearing: An approach to distributed, eventually consistent edge computation. In *2015 IEEE 34th Symposium on Reliable Distributed Systems Workshop (SRDSW)*. IEEE, 62–67.
- [27] Christopher S. Meiklejohn, Vitor Enes, Junghun Yoo, Carlos Baquero, Peter Van Roy, and Annette Bieniusa. 2017. Practical Evaluation of the Lasp Programming Model at Large Scale: An Experience Report. In *Proceedings of the 19th International Symposium on Principles and Practice of Declarative Programming (PPDP '17)*. ACM, New York, NY, USA, 109–114. <https://doi.org/10.1145/3131851.3131862>
- [28] Aljoscha Meyer. 2018. *SSB Specification*. <https://spec.scuttlebutt.nz/feed/messages.html>
- [29] Ilya Moiseenko and Lixia Zhang. 2014. Consumer-producer API for Named Data Networking. In *Proceedings of the 1st ACM Conference on Information-Centric Networking (ACM-ICN '14)*. ACM, New York, NY, USA, 177–178. <https://doi.org/10.1145/2660129.2660158>
- [30] Philipp Moll, Sebastian Theuermann, and Hermann Hellwagner. 2018. Persistent Interests in Named Data Networking. In *2018 IEEE 87th Vehicular Technology*

- Conference (VTC Spring)*. IEEE, 1–5.
- [31] David Nelson. 2011. *Crypto-Agility Requirements for Remote Authentication Dial-In User Service (RADIUS)*. Technical Report.
 - [32] Noffle. 2016. *git-ssb-intro*. <https://github.com/noffle/git-ssb-intro#push-conflicts>
 - [33] Diego Ongaro and John K. Ousterhout. 2014. In Search of an Understandable Consensus Algorithm. In *Proc USENIX Annual Technical Conference*. 305–319. <https://www.usenix.org/system/files/conference/atc14/atc14-paper-ongaro.pdf>
 - [34] Mikko Pitkänen, Teemu Kärkkäinen, Jörg Ott, Marco Conti, Andrea Passarella, Silvia Giordano, Daniele Puccinelli, Franck Legendre, Sacha Trifunovic, Karin Hummel, et al. 2012. SCAMPI: Service platform for social aware mobile and pervasive computing. In *Proceedings of the first edition of the MCC workshop on Mobile cloud computing*. ACM, 7–12.
 - [35] Mauro Sardara, Luca Muscariello, and Alberto Compagno. 2018. A Transport Layer and Socket API for (H)ICN: Design, Implementation and Performance Analysis. In *Proceedings of the 5th ACM Conference on Information-Centric Networking (ICN '18)*. ACM, New York, NY, USA, 137–147. <https://doi.org/10.1145/3267955.3267972>
 - [36] Bruce Schneier and John Kelsey. 1998. Cryptographic support for secure logs on untrusted machines.. In *USENIX Security Symposium*, Vol. 98. 53–62.
 - [37] James Scott, Jon Crowcroft, Pan Hui, and Christophe Diot. 2006. Haggle: A networking architecture designed around mobile users.
 - [38] Marc Shapiro, Nuno Preguiça, Carlos Baquero, and Marek Zawirski. 2011. Conflict-free replicated data types. In *Symposium on Self-Stabilizing Systems*. Springer, 386–400.
 - [39] Atul Singh et al. 2006. Eclipse attacks on overlay networks: Threats and defenses. In *IEEE INFOCOM*. Citeseer.
 - [40] Dominic Tarr, Paul Frazee, Christian Bundy, Matt McKegg, Anders Rune Jensen, Mix Irving, et al. 2014. *SSB Server*. <https://github.com/ssbc/ssb-server>
 - [41] Dominic Tarr, Mix Irving, Christian Bundy, Michael Williams, Anders Rune Jensen, Andre Staltz, and Matt McKegg. 2014. *SSB Server*. <https://github.com/ssbc/ssb-friends>
 - [42] Christian Tschudin. 2019. A Broadcast-Only Communication Model Based on Replicated Append-Only Logs. *SIGCOMM Comput. Commun. Rev.* 49, 2 (May 2019), 37–43. <https://doi.org/10.1145/3336937.3336943>
 - [43] Yingdi Yu, Alexander Afanasyev, Jan Seedorf, Zhiyi Zhang, and Lixia Zhang. 2017. NDN DeLorean: An authentication system for data archives in named data networking. In *Proceedings of the 4th ACM Conference on Information-Centric Networking*. ACM, 11–21.
 - [44] Zhenkai Zhu, Chaoyi Bian, Alexander Afanasyev, Van Jacobson, and Lixia Zhang. 2012. Chronos: Serverless multi-user chat over NDN. *Technical Report NDN-0008* (2012).